OpenMP
C/C++

OpenMP 3.1 API C/C++ Syntax Quick Reference Card

OpenMP Application Program Interface (API) is a portable, scalable model that gives shared-memory parallel programmers a simple and flexible interface for developing parallel applications for platforms ranging from the desktop to the supercomputer.

OpenMP supports multi-platform shared-memory parallel programming in C/C++ and Fortran on all architectures, including Unix platforms and Windows NT platforms.

A separate OpenMP reference card for Fortran is also available.

[n.n.n] refers to sections in the OpenMP API Specification available at www.openmp.org.

Directives

An OpenMP executable directive applies to the succeeding structured block or an OpenMP Construct. A *structured-block* is a single statement or a compound statement with a single entry at the top and a single exit at the bottom.

Parallel [2.4]
The **parallel** construct forms a team of threads and starts parallel execution.

#pragma omp parallel *[clause[[,]clause] ...]*
　　structured-block

clause:
　　　if(*scalar-expression*)
　　　num_threads(*integer-expression*)
　　　default(**shared** | **none**)
　　　private(*list*)
　　　firstprivate(*list*)
　　　shared(*list*)
　　　copyin(*list*)
　　　reduction(*operator*: *list*)

Loop [2.5.1]
The **loop** construct specifies that the iterations of loops will be distributed among and executed by the encountering team of threads.

#pragma omp for *[clause[[,]clause] ...]*
　　for-loops

clause:
　　　private(*list*)
　　　firstprivate(*list*)
　　　lastprivate(*list*)
　　　reduction(*operator*: *list*)
　　　schedule(*kind[, chunk_size]*)
　　　collapse(*n*)
　　　ordered
　　　nowait

> **Most common form of the for loop:**
>
> for(var = lb;
>   var relational-op b;
>   var += incr)

kind:
- **static**: Iterations are divided into chunks of size *chunk_size*. Chunks are assigned to threads in the team in round-robin fashion in order of thread number.
- **dynamic**: Each thread executes a chunk of iterations then requests another chunk until no chunks remain to be distributed.
- **guided**: Each thread executes a chunk of iterations then requests another chunk until no chunks remain to be assigned. The chunk sizes start large and shrink to the indicated *chunk_size* as chunks are scheduled.
- **auto**: The decision regarding scheduling is delegated to the compiler and/or runtime system.
- **runtime**: The schedule and chunk size are taken from the *run-sched-var* ICV.

Sections [2.5.2]
The **sections** construct contains a set of structured blocks that are to be distributed among and executed by the encountering team of threads.

#pragma omp sections *[clause[[,] clause] ...]*
```
   {
   [#pragma omp section]
      structured-block
   [#pragma omp section
      structured-block]
   ...
   }
```
clause:
　　　private(*list*)
　　　firstprivate(*list*)
　　　lastprivate(*list*)
　　　reduction(*operator*: *list*)
　　　nowait

Single [2.5.3]
The **single** construct specifies that the associated structured block is executed by only one of the threads in the team (not necessarily the master thread), in the context of its implicit task.

#pragma omp single *[clause[[,]clause] ...]*
　　structured-block

clause:
　　　private(*list*)
　　　firstprivate(*list*)
　　　copyprivate(*list*)
　　　nowait

Parallel Loop [2.6.1]
The parallel loop construct is a shortcut for specifying a **parallel** construct containing one or more associated loops and no other statements.

#pragma omp parallel for *[clause[[,]clause] ...]*
　　for-loop

clause:
　　Any accepted by the **parallel** or **for** directives, except the **nowait** clause, with identical meanings and restrictions.

> **Simple Parallel Loop Example**
> The following example demonstrates how to parallelize a simple loop using the parallel loop construct.
>
> ```
> void simple(int n, float *a, float *b)
> {
>    int i;
> #pragma omp parallel for
>    for (i=1; i<n; i++) /* i is private by default */
>       b[i] = (a[i] + a[i-1]) / 2.0;
> }
> ```

Parallel Sections [2.6.2]
The **parallel sections** construct is a shortcut for specifying a **parallel** construct containing one **sections** construct and no other statements.

#pragma omp parallel sections *[clause[[,]clause] ...]*
```
   {
   [#pragma omp section]
      structured-block
   [#pragma omp section
      structured-block
   ...
   }
```
clause:
　　Any of the clauses accepted by the **parallel** or **sections** directives, except the **nowait** clause, with identical meanings and restrictions.

Task [2.7.1]
The **task** construct defines an explicit task. The data environment of the task is created according to the data-sharing attribute clauses on the **task** construct and any defaults that apply.

#pragma omp task *[clause[[,]clause] ...]*
　　structured-block

clause:
　　　if(*scalar-expression*)
　　　final(*scalar-expression*)
　　　untied
　　　default(**shared** | **none**)
　　　mergeable
　　　private(*list*)
　　　firstprivate(*list*)
　　　shared(*list*)

Taskyield [2.7.2]
The **taskyield** construct specifies that the current task can be suspended in favor of execution of a different task.

#pragma omp taskyield

Master [2.8.1]
The **master** construct specifies a structured block that is executed by the master thread of the team. There is no implied barrier either on entry to, or exit from, the **master** construct.

#pragma omp master
　　structured-block

Critical [2.8.2]
The **critical** construct restricts execution of the associated structured block to a single thread at a time.

#pragma omp critical *[(name)]*
　　structured-block

Barrier [2.8.3]
The **barrier** construct specifies an explicit barrier at the point at which the construct appears.

#pragma omp barrier

Taskwait [2.8.4]
The **taskwait** construct specifies a wait on the completion of child tasks of the current task.

#pragma omp taskwait

Atomic [2.8.5]
The **atomic** construct ensures that a specific storage location is updated atomically, rather than exposing it to the possibility of multiple, simultaneous writing threads.

#pragma omp atomic *[read | write | update | capture]*
　　expression-stmt

#pragma omp atomic capture
　　structured-block

where *expression-stmt* may be one of the following forms:

if *clause* is...	*expression-stmt*:		
read	v = x;		
write	x = expr;		
update or is not present	x++;　　　　　x--; --x;	x binop = expr;	++x; x = x binop expr;
capture	v = x++;　　　v = x--; v = --x;	v = x binop= expr;	v = ++x;

and *structured-block* may be one of the following forms:
{v = x; x binop= expr;}　　　{x binop= expr; v = x;}
{v = x; x = x binop expr;}　　{x = x binop expr; v = x;}
{v = x; x++;}　　{v = x; ++x;}　　{++x; v = x;}　　{x++; v = x;}
{v = x; x--;}　　{v = x; --x;}　　{--x; v = x;}　　{x--; v = x;}

Flush [2.8.6]
The **flush** construct executes the OpenMP flush operation, which makes a thread's temporary view of memory consistent with memory, and enforces an order on the memory operations of the variables.

#pragma omp flush *[(list)]*

Ordered [2.8.7]
The **ordered** construct specifies a structured block in a loop region that will be executed in the order of the loop iterations. This sequentializes and orders the code within an **ordered** region while allowing code outside the region to run in parallel.

#pragma omp ordered
　　structured-block

Threadprivate [2.9.2]
The **threadprivate** directive specifies that variables are replicated, with each thread having its own copy.

#pragma omp threadprivate(*list*)
list:
　　A comma-separated list of file-scope, namespace-scope, or static block-scope variables that do not have incomplete types.

Runtime Library Routines

Execution Environment Routines [3.2]
Execution environment routines affect and monitor threads, processors, and the parallel environment.

void omp_set_num_threads(
    int *num_threads*);
Affects the number of threads used for subsequent **parallel** regions that do not specify a **num_threads** clause.

int omp_get_num_threads(void);
Returns the number of threads in the current team.

int omp_get_max_threads(void);
Returns maximum number of threads that could be used to form a new team using a **parallel** construct without a **num_threads** clause.

int omp_get_thread_num(void);
Returns the ID of the encountering thread where ID ranges from zero to the size of the team minus 1.

int omp_get_num_procs(void);
Returns the number of processors available to the program.

Data Types For Runtime Library Routines

omp_lock_t: Represents a simple lock.

omp_nest_lock_t: Represents a nestable lock.

omp_sched_t: Represents a schedule.

int omp_in_parallel(void);
Returns **true** if the call to the routine is enclosed by an active **parallel** region; otherwise, it returns **false**.

void omp_set_dynamic(
    int *dynamic_threads*);
Enables or disables dynamic adjustment of the number of threads available for the execution of **parallel** regions by setting the value of the *dyn-var* ICV.

int omp_get_dynamic(void);
Returns the value of the *dyn-var* ICV, determining whether dynamic adjustment of the number of threads is enabled or disabled.

void omp_set_nested(int *nested*);
Enables or disables nested parallelism, by setting the *nest-var* ICV.

int omp_get_nested(void);
Returns the value of the *nest-var* ICV, which determines if nested parallelism is enabled or disabled.

void omp_set_schedule(omp_sched_t *kind*,
    int *modifier*);
Affects the schedule that is applied when **runtime** is used as schedule kind, by setting the value of the *run-sched-var* ICV.

kind is one of **static**, **dynamic**, **guided**, **auto**, or an implementation-defined schedule. See **loop** construct [2.5.1] for descriptions.

void omp_get_schedule(
    omp_sched_t **kind*, int **modifier*);
Returns the value of *run-sched-var* ICV, which is the schedule applied when **runtime** schedule is used.

See *kind* above.

int omp_get_thread_limit(void);
Returns the value of the *thread-limit-var* ICV, which is the maximum number of OpenMP threads available to the program.

void omp_set_max_active_levels(
    int *max_levels*);
Limits the number of nested active **parallel** regions, by setting *max-active-levels-var* ICV.

int omp_get_max_active_levels(void);
Returns the value of *max-active-levels-var* ICV, which determines the maximum number of nested active **parallel** regions.

int omp_get_level(void);
Returns the number of nested **parallel** regions enclosing the task that contains the call.

int omp_get_ancestor_thread_num(
    int *level*);
Returns, for a given nested level of the current thread, the thread number of the ancestor or the current thread.

int omp_get_team_size(int *level*);
Returns, for a given nested level of the current thread, the size of the thread team to which the ancestor or the current thread belongs.

int omp_get_active_level(void);
Returns the number of nested, active **parallel** regions enclosing the task that contains the call.

int omp_in_final(void);
Returns **true** if the routine is executed in a **final** or included task region; otherwise, it returns **false**.

Lock Routines [3.3]
Lock routines support synchronization with OpenMP locks.

void omp_init_lock(omp_lock_t **lock*);
void omp_init_nest_lock(
    omp_nest_lock_t **lock*);
These routines initialize an OpenMP lock.

void omp_destroy_lock(omp_lock_t **lock*);
void omp_destroy_nest_lock(
    omp_nest_lock_t **lock*);
These routines ensure that the OpenMP lock is uninitialized.

void omp_set_lock(omp_lock_t **lock*);
void omp_set_nest_lock(
    omp_nest_lock_t **lock*);
These routines provide a means of setting an OpenMP lock.

void omp_unset_lock(omp_lock_t **lock*);
void omp_unset_nest_lock(
    omp_nest_lock_t **lock*);
These routines provide a means of unsetting an OpenMP lock.

int omp_test_lock(omp_lock_t **lock*);
int omp_test_nest_lock(
    omp_nest_lock_t **lock*);
These routines attempt to set an OpenMP lock but do not suspend execution of the task executing the routine.

Timing Routines [3.4]
Timing routines support a portable wall clock timer.

double omp_get_wtime(void);
Returns elapsed wall clock time in seconds.

double omp_get_wtick(void);
Returns the precision of the timer used by **omp_get_wtime**.

Clauses

The set of clauses that is valid on a particular directive is described with the directive. Most clauses accept a comma-separated list of list items. All list items appearing in a clause must be visible.

Data Sharing Attribute Clauses [2.9.3]
Data-sharing attribute clauses apply only to variables whose names are visible in the construct on which the clause appears.

default(shared | none)
Controls the default data-sharing attributes of variables that are referenced in a **parallel** or **task** construct.

shared(*list*)
Declares one or more list items to be shared by tasks generated by a **parallel** or **task** construct.

private(*list*)
Declares one or more list items to be private to a task.

firstprivate(*list*)
Declares one or more list items to be private to a task, and initializes each of them with the value that the corresponding original item has when the construct is encountered.

lastprivate(*list*)
Declares one or more list items to be private to an implicit task, and causes the corresponding original item to be updated after the end of the region.

reduction(*operator***:***list*)
Declares accumulation into the list items using the indicated associative operator. Accumulation occurs into a private copy for each list item which is then combined with the original item.

Operators for reduction (initialization values)			
+	(0)	|	(0)
*	(1)	^	(0)
-	(0)	&&	(1)
&	(~0)	||	(0)
max (Least number in **reduction** list item type)			
min (Largest number in **reduction** list item type)			

Data Copying Clauses [2.9.4]
These clauses support the copying of data values from private or threadprivate variables on one implicit task or thread to the corresponding variables on other implicit tasks or threads in the team.

copyin(*list*)
Copies the value of the master thread's threadprivate variable to the threadprivate variable of each other member of the team executing the **parallel** region.

copyprivate(*list*)
Broadcasts a value from the data environment of one implicit task to the data environments of the other implicit tasks belonging to the **parallel** region.

Environment Variables

Environment variables are described in section [4] of the API specification. Environment variable names are upper case, and the values assigned to them are case insensitive and may have leading and trailing white space.

OMP_SCHEDULE *type[,chunk]*
Sets the *run-sched-var* ICV for the runtime schedule type and chunk size. Valid OpenMP schedule types are **static**, **dynamic**, **guided**, or **auto**. *chunk* is a positive integer that specifies chunk size.

OMP_NUM_THREADS *list*
Sets the *nthreads-var* ICV for the number of threads to use for **parallel** regions.

OMP_DYNAMIC *dynamic*
Sets the *dyn-var* ICV for the dynamic adjustment of threads to use for **parallel** regions. Valid values for *dynamic* are **true** or **false**.

OMP_PROC_BIND *bind*
Sets the value of the global *bind-var* ICV. The value of this environment variable must be **true** or **false**.

OMP_NESTED *nested*
Sets the *nest-var* ICV to enable or to disable nested parallelism. Valid values for *nested* are **true** or **false**.

OMP_STACKSIZE *size[B | K | M | G]*
Sets the *stacksize-var* ICV that specifies the size of the stack for threads created by the OpenMP implementation. *size* is a positive integer that specifies stack size. If unit is not specified, *size* is measured in kilobytes (K).

OMP_WAIT_POLICY *policy*
Sets the *wait-policy-var* ICV that controls the desired behavior of waiting threads. Valid values for *policy* are **ACTIVE** (waiting threads consume processor cycles while waiting) and **PASSIVE**.

OMP_MAX_ACTIVE_LEVELS *levels*
Sets the *max-active-levels-var* ICV that controls the maximum number of nested active **parallel** regions.

OMP_THREAD_LIMIT *limit*
Sets the *thread-limit-var* ICV that controls the maximum number of threads participating in the OpenMP program.